

Final Report: June 10

Lecturer: Eric Darve

Student: Nurbek Tazhimbetov

1.1 GEMM and Grading Mode 4

First, I tested the implementation of myGEMM using grading mode 4,

	GEMM1	GEMM2
Serial Time	0.0468626	0.00441485
Vanila Parallel	0.262259	0.00698588
Shared Parallel	0.15199	0.0065526
Parallel Shared A Plus	0.0753727	0.00548108
Relative difference	1.85746e-16	3.2505e-16

As you can see for GEMM1, I have started from vanilla parallel implementation and got around 0.26 seconds, while after implementing shared version, I nearly doubled the speed with 0.15 seconds. In the last algorithm, we nearly quadrupled the initial time and achieved 0.075 seconds. However, we are still far away from the serial code, which is 0.047 seconds. The shared A plus algorithm has been described in *Benchmarking GPUs to Tune Dense Linear Algebra* by Demmel and Volkov in 2008.

1.2 Grading Modes 3, 2 & 1

Table 1.1: Results for grading mode 3, 2 & 1

	Grading mode 3	Grading mode 2	Grading mode 1	./main
Serial	13.5539 s	50.2858 s	196.722 s	
Parallel	12.4649 s	13.6078 s	52.6221 s	107.304 s
Cross Validation accuracy (both)	0.845167	0.924167	0.894333	0.909833
Serial New Time	10.9739 s	31.4504 s	125.92 s	
Parallel New Time	9.13504 s	4.13525 s	14.2701 s	37.5473 s
Cross Validation New accuracy	0.845167	0.924167	0.894333	0.909833
$\ W_1^{(s)} - W_1^{(p)}\ _{\max}$	4.24529e-11	4.57298e-09	1.23397e-15	
$\ W_2^{(s)} - W_2^{(p)}\ _{\max}$	2.0301e-13	3.46816e-11	4.27981e-16	
$\ b_1^{(s)} - b_1^{(p)}\ _{\max}$	1.0906e-10	6.71967e-09	2.99972e-15	
$\ b_2^{(s)} - b_2^{(p)}\ _{\max}$	2.57021e-12	3.29032e-10	6.13302e-16	
$\ W_1^{(s)} - W_1^{(p)}\ _2$	4.04886e-11	3.47396e-09	1.09988e-15	
$\ W_2^{(s)} - W_2^{(p)}\ _2$	8.70559e-13	6.0723e-11	7.15398e-16	
$\ b_1^{(s)} - b_1^{(p)}\ _2$	2.98602e-11	2.08037e-09	1.50882e-15	
$\ b_2^{(s)} - b_2^{(p)}\ _2$	2.83175e-12	2.61242e-10	5.87586e-16	

Next, I used grading mode 3, 2 and 1 to run the neural net training and got the following results. As you can see, the time has dropped significantly. This has been achieved by copying and scattering the data outside of the epoch loop. We call `MPI_Scatterv()` only for one epoch for each batch iteration and store them on the device. This may seem unwise memory-wise, however, accessing to the stored X and y values is much faster from the device itself rather copying from the host every time.

1.3 Profiling

Table 1.2: Results from .err file for ./main

Old Time (%)	Old Time	# of Old Calls	Old Name	New Time (%)	New Time	# of New Calls	New Name
36.59%	19.1606s	19040	[CUDA memcpy HtoD]	31.51%	5.68019s	5442	[CUDA memcpy HtoD]
21.98%	11.5084s	2720	GEMMTransposeKernel	28.94%	5.21721s	5440	[CUDA memcpy DtoH]
20.41%	10.6877s	2720	GEMMKernel	16.21%	2.92294s	2720	GEMMTransposeKernel
19.14%	10.0203s	10880	[CUDA memcpy DtoH]	15.83%	2.85349s	1360	GEMMSigmoidKernel
0.65%	342.03ms	5440	GradientKernel	4.23%	763.00ms	1360	GEMMAdditionKernel
0.59%	309.73ms	1360	HadamardKernel	1.72%	310.11ms	1360	HadamardKernel
0.46%	239.29ms	2720	SumOfRowKernel	1.43%	257.61ms	2720	SumOfRowKernel
0.14%	72.683ms	1360	SigmoidKernel	0.08%	13.700ms	1360	SumOfExpColKernel
0.03%	13.689ms	1360	SumOfExpColKernel	0.04%	6.6954ms	1360	SoftmaxKernel
0.01%	6.7247ms	1360	SoftmaxKernel	0.02%	4.0574ms	1360	[CUDA memcpy DtoD]

Another interesting topic to discuss is the difference between serial and parallel results for W_1, W_2, b_1 and b_2 . This can be explained with the fact that some of the numerical computations are done on the GPU

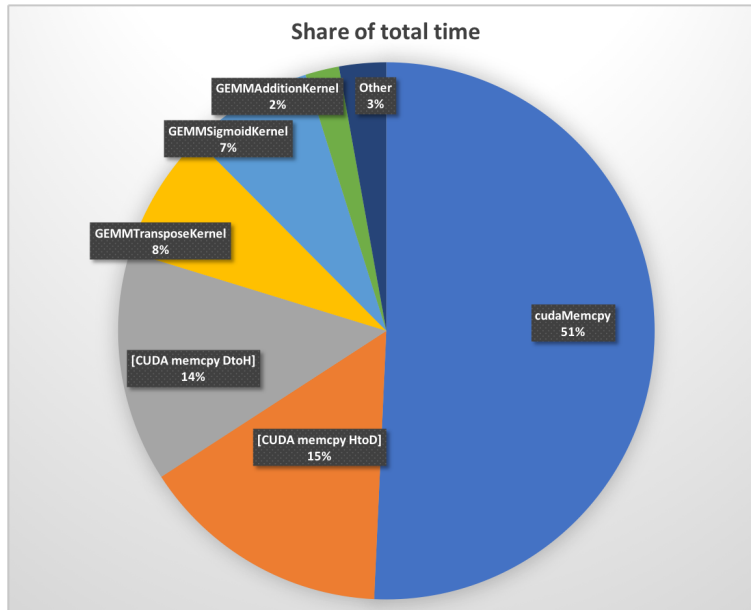


Figure 1.1: Share of total time

and CPU for parallel and serial algorithms, respectively. Therefore, even on hardware level we might have differences between these codes. The last interesting issue here is that for the grading mode 2, the error increases and is higher than for the other two. Perhaps, this is because the learning rate of this algorithm and number of iterations are not a good match.

Table 1.3: Results from .err file for ./main

%	Old Time	# of Old Calls	Old Name	%	New Time	# of New Calls	New Name
98.81	54.8160s	29920	cudaMemcpy	96.75	19.0458s	12242	cudaMemcpy
0.71	394.53ms	13	cudaMalloc	2.16	425.68ms	11	cudaMalloc
0.34	187.74ms	19040	cudaLaunch	0.72	141.88ms	12240	cudaLaunch
0.05	29.384ms	364	cuDeviceGetAttribute	0.14	27.227ms	364	cuDeviceGetAttribute
0.04	22.668ms	104720	cudaSetupArgument	0.10	19.373ms	77520	cudaSetupArgument
0.02	12.384ms	13	cudaFree	0.07	13.786ms	11	cudaFree
0.02	8.4327ms	19040	cudaConfigureCall	0.04	6.9613ms	12240	cudaConfigureCall
0.01	2.8570ms	4	cuDeviceGetName	0.02	3.0728ms	4	cuDeviceTotalMem
0.00	2.7006ms	4	cuDeviceTotalMem	0.01	2.4341ms	4	cuDeviceGetName
0.00	19.218us	1	cudaSetDevice	0.00	20.598us	1	cudaSetDevice
0.00	4.8910us	3	cuDeviceGetCount	0.00	16.675us	2	cudaFreeHost
0.00	4.6690us	12	cuDeviceGet	0.00	6.1060us	12	cuDeviceGet
0.00	2.8930us	1	cudaGetDeviceCount	0.00	4.7760us	3	cuDeviceGetCount
				0.00	2.9900us	1	cudaGetDeviceCount

1.4 Optimization




Figure 1.2: Visual Profiler Results

As it was expected the most time-consuming action is `memcpy`; both from the device to the host and vice-versa. Since the A Plus algorithm has been implemented, we have achieved (as mentioned before) four times faster solver; `GEMMTransposeKernel` dropped from over 11 seconds to just under 3 seconds and `GEMMKernel` dropped from over 10 seconds to almost 3 seconds. Furthermore, since we are storing X and y on the device, I minimized the number of calls of the copying function between the host and device (`memcpy`) from 29,920 to 12,242. This saved me 35 seconds.

Same can be said about the GEMM. As we have discussed earlier, the GEMM's fastest implementation is almost four times faster than the vanilla GEMM. That gave me almost 17 seconds.

1.5 Possible Further Optimization

After reducing the values of W_1 , W_2 , b_1 and b_2 on the host, I have used `arma`'s matrix-matrix addition instead of copying back to the device and write my own matrix-matrix addition, although for the preliminary test I had the GPU version of the gradient descent update. I tend to think this is a minor improvement, however, it might be important, if we have a bigger batch size, which increases number of computations. I could have also minimized the number of kernels, however, as you can see from Figure 1.2, the other kernels are not very time-consuming. However, this is still might give us more time.

- The percentage of time when `memcpy` is being performed in parallel with kernel is low. Therefore, we might want to copy in parallel.
- The percentage of of time when two kernels are being executed in parallel is low. In the back propagation part, we might actually be able to get more time.